

Name: HyunJong Lee

GTID: 902987836

Email: hlee652@gatech.edu

Class: CS4210 - A

Final Report for GTThreads

GTThreads package:

O(1) scheduler:

GT thread package implements O(1) scheduler for insertion, removal and retrieval of next uthread to be scheduled from the queue. For each of kthread, which acts like a virtual CPU, has two runqueues: active and expired respectively. When runqueue is initialized, it puts all the tasks into active runqueue. When each task is preempted, it puts into expired runqueue. Once the active runqueue becomes empty, it swaps the active and expired runqueue. Swapping two runqueues is done by swapping two pointers and can be done in O(1) time.

As provided in the GTThread better understanding pdf file, sched_find_best_thread() performs following work:

1. Try to find the highest priority RUNNABLE uthread in active runqueue by setting the lowest bit in the uthread_mask to gain the highest priority index.
2. Using the priority index, it takes priority queue and using the group_mask, it takes uthread group.
3. Push the corresponding uthread to the tail of the queue and removes from the active runqueue

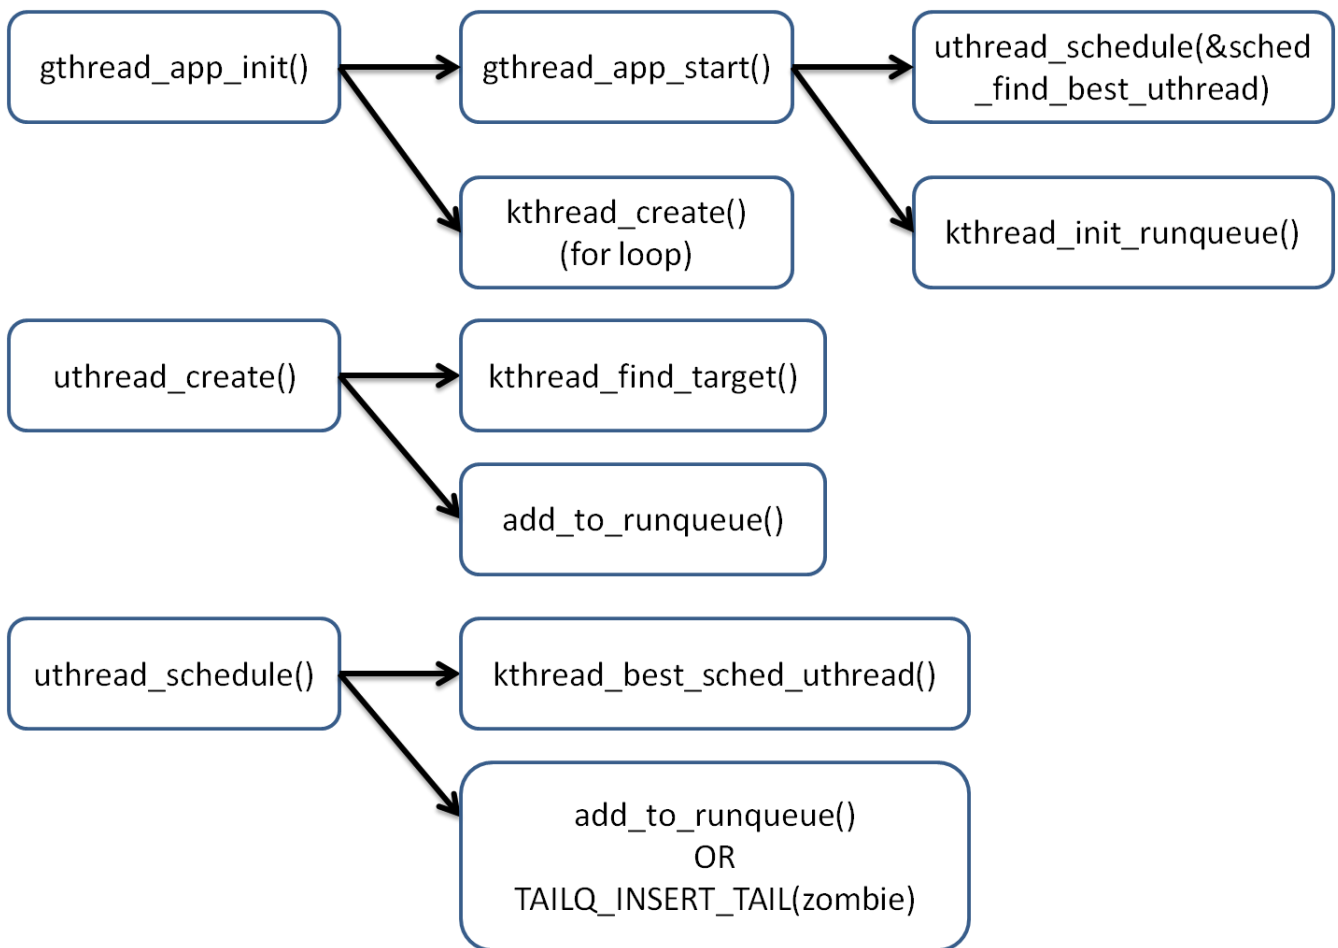


Fig1. Diagram for basic function calls in gthreads.

Credit-based scheduler:

How it works:

Credit scheduler works somewhat similar way as priority queue scheduler works. At instantiation state, it assigns default credit to each of uthreads. Given the credit, each uthread runs for given amount of time, calculated based on credit in milliseconds. Once the uthread is scheduled, it loses certain amount of credit, based on time it spent on CPU. Until the credit remains in positive number, it resides in active runqueue. However, as soon as it consumes all the credit, it is moved to expired runqueue with newly assigned credit (default credit).

Implementation in GTThreads:

Each of uthread is assigned default credit and credit left in uthread_create() and added to active runqueue:

```

u_new->credits.credit_left = u_new->credits.def_credit = credits;
add_to_runqueue(kthread_runq->active_runq, &(kthread_runq->kthread_runqlock), u_new);

```

When VTALRM signal is received by one of the kthreads, the handler function ksched_priority() is invoked to schedule a new uthread by calling uthread_schedule() function.

uthread_schedule() function is where credit scheduler algorithm is implemented. Within the function, it

calls `calculate()`, which updates the time it is updated and credit left, which is decreased by current time – last updated time in millisecond.

If the credit is less than 0, it adds to expired runqueue. Otherwise, it adds to active runqueue. Before calling the `setlongjmp()`, it sets the timer. If the `uthread` is first time scheduled (`== def_credit`), it uses `kthread_init_vtalm_timeslice()` to set timer. Otherwise, it converts remaining credits into millisecond and set the timer.

For the load balancing part, it briefly checks whether it's first time scheduled thread. If so, it goes through `kthreads` array and add into specific runqueue.

To cover `gt_yield()` API call, I added variable `yid`, to check voluntarily preempted case. Note that, the real work (computation), is done in `uthread_context_func()`.

Summarization of results for all the test cases:

The following chart shows average and standard deviation for both running time (on CPU) and total execution time for each of unique pair (metrics size and credit).

Metrics Credit	Avg. Running time	Avg. Execution time	Standard deviation Running time	Standard deviation Execution time
(32x 32 25)	19	1107	2	121
(32x 32 50)	17	1080	1	122
(32x 32 75)	17	827	1	451
(32x 32 100)	17	237	1	427
(64x 64 25)	133	1563	13	127
(64x 64 50)	130	1427	11	126
(64x 64 75)	131	1391	9	88
(64x 64 100)	134	1252	2	77
(128x128 25)	1043	4801	845	1313
(128x128 50)	1024	4191	103	909
(128x128 75)	1027	3921	1409	520
(128x128 100)	1050	3027	77	839

(256x256 25)	9835	32744	1880	3018
(256x256 50)	8915	24977	362	2856
(256x256 75)	8979	24907	130	2973
(256x256 100)	8682	15985	771	2789

Regardless of amount of credit given, it shows similar running time on CPU for same metrics size. That is to say, even though the thread with higher credit finishes the job early, at the end, it takes same amount of time in CPU on doing its job. A notable difference is average total execution time. As we expected, the higher credit it has, the lower waiting time it takes. Since a thread with higher credit gets more timeslice in given order, it can reside in CPU for longer time and finish the job early. Therefore, it waits less to load itself onto the CPU.

Implementation Issues:

Despite the uthread in general is implemented sufficiently, it outputs illegal instruction or segmentation fault frequently. In such case, you can run multiple times until you get the statistics.